9 month review report

Candidate : Bokar N'Diaye

Supervisors : Tobias Blanke, Paula Helm, Delfina Sol Martinez Pandiani

Thesis

The four proposed papers all aim to draw attention to the *situatedness* of actors involved in the controversies that surround production and adoption of state-of-the-art deep learning models "in the wild" – especially in the context of subcultural scenes that claim to use and understand them *oppositionally* to perceived mainstream norms. We will aim to "explain and explicate" (in the sense of Clifford Geertz) the reactions of early adoption or vehement rejection of these models, in the face of their possible *opacity*, as well as imperatives of *secrecy* that the different actors (users, opponents, experts, companies or regulator) aim to strategically maintain. We hope of establishing insights, especially by showing that folk ontologies, heuristics or *bricolage* methods (from practices of vibe-coding to jailbreaking or "slop" detection) are an integral part of the plurality of digital literacies, and that paying attention to the processes of sense-making, world-making and collective organization that occur in these "ecosystems of thoughts" may provide crucial context towards understanding and communicating about human/machine interactions in the current era of deep learning.

Scientific Contribution

- **"Tech Oligarchy, Digital Infrastructure, and the Politics of Data Stifling"**, a collective paper co-written with Louis Ravn and submitted to the Science as Culture Forum, a first look at a top-down "prosopography" of elite tech actors.
- **"Take a S-EAT, Eat Some TREACLES : Mapping the Middlegrounds and Middlemen of AI on the LessWrong forum",** A draft for the first paper of my thesis (described below).

Paper 1 : Epistemic Standards of the LessWrong

Timeline : Year 1

Objective and Methodology: *(presented in more detail in joint draft)*

The article would propose to reformulate the concept of the TESCREAL bundle as top-down explanation of AI controversies, instead underlining the role of specific digital "social worlds" in the creation and circulation of such narratives. It provides a retrospective, contextualized outlook on the critical role played by an oppositional or "cultic milieu" (Campbell 1972) in the production and adoption of generative models in the 2020s, and evidences their impact on the membership, terminology, and assumptions of an emerging and rapidly unfolding field of research and industrial development (explainable / interpretable AI). Lastly, it would underline the role of ML agents as both ostensive devices and co-authors of the community's staged efforts of sense-making, and critically mobilize their terminology as a pluralizing lens to genAI controversies and situations.

Target Publications: First Monday, Sociologica

Comments by Delfina:
- as someone who doesnt know LessWrong, more context is needed to be convincing of the role that these platforms have, and that they are not just very insular spaces. Is there any work tracing how stuff from LessWrong travels out to other (digital and material) platforms?

- glitch tokens seems a very interesting case study to understand the epistemic standards and commitments evolving on LessWrong

Paper 2 : Questioning Prominence : Comparing Collective Memory in LLMs and the Pantheon Database

Timeline : Year 1-2

Objective: This paper seeks to replicate and nuance experiments such as Nicolas Kayser-Bril's on the apparent convergence of large language models when prompted to constitute a ranking of the 200. Kayser-Bril concludes that LLMs share a homogenized worldview that reflects the norms of its Californian developers. While persuasive, this claim also relies primarily on what Fabian Offert and Ranjodh Singh Dhaliwal term "stack casuistry" : the affirmation of a LLM, without due untangling of the emergent factors of its computation or. To provide a commensurable measure of prominence grounded in documented collective memory, the study will compare LLM-generated predicted listings of "the most important historical figures" with the similar rankings established in the Pantheon 2.0 ((Yu et al., 2016) and Networked Pantheon (Beytía & Schobin, 2020) datasets, based on Wikipedia metrics. The goal is to assess how far model-based hierarchies of fame diverge from empirically observed human attention over time, geography, and occupation, and to determine whether closer attention to the probabilistic workings of the model sheds light on the asymmetries and anomalies of its output.

Relevance : The project aligns with the current turn toward explainable and interpretable AI by employing interpretable probabilistic indicators—namely, token-level logprobabilities—as a means to inspect the internal hierarchies implicit in autoregressive text generation. It will contrast the evolving logprob structures of LLMs with Wikipedia's collectively curated architectures of visibility, treating both as socio-technical mirrors of global knowledge production. Attentive to their blind spots, the paper will also explore opacity and apophasis—the figures, domains, and civilizations that LLMs systematically omit or marginalize—thus relating technical bias measurement to cultural hermeneutics.

Methodology : Several leading language models (GPT-4, Claude-3, Gemini, Mistral, etc.) will be prompted to produce ranked lists of the "n-100 most important figures in history," both globally and within specific occupational or temporal subsets. Statistical correlations, divergence metrics, and network analyses will then be used to map under- and over-representation patterns relative to the Pantheon database.

Target Publications : Weizenbaum Journal of the Digital Society

I am also interested in this: specifically investigating these phenomena, potentially through the lens of 'queer failure' (Halberstam)

Paper 3 : Abstention, Apophasis and Agnotology

Timeline : Year 2-3

Objective: The contribution aims to evidence the ability of the model to draw attention to the things it cannot (or refuses to) mention, despite being expected to. The project investigates how a model's silences, refusals, and errors can be more instructive than its correct outputs. We extend this reflection into three key phenomena: *abstention* (the engineered or emergent mechanisms through which a model withholds information), *apophasis* (the rhetorical and psychological effect where the model's act of refusal makes the unmentioned topic conspicuously present and significant), and *agnotology* (the study of how becomes a rich site for understanding the model's architecture, biases, and worldview, as well as the cultural context and rationales that shape it.) Taking inspiration from Anders' Munk "thick machines", a proposition to consider LLMs as able to provide anthropological insight through their imitation communicational patterns in a studied community, the project seeks to design and analyze AI systems that, by failing in specific and meaningful ways, actively teach users to value these limitations as potential advantages of these models.

Methodology : Mixed-methods ethnography + model experimentation. Based on Gabriele de Seta's practices of self-teaching vernacular AI as source of ethnographic insights, the article would document the constitution of three projects would serve as field devices, incorporating a measure of perversity and complicating the idea of seamless "interactivity". Each reflective of a "social media ritual" indexable through the collection of multimodal corpora, the models are exploited as "thick machines", deriving cultural context through imperfect participation.

- **A visual model**, based on the practices of a long-established internet forum practice "photoshop requests". The activity involves specialists either genuinely fulfilling an image-editing request or interpreting it with deliberately literal, "malicious" humor (e.g., a request to "make me look an inch closer to my friend" resulting in the user's face being photoshopped on a finger). An open-source multimodal model able to enact topical changes to an image would be trained on requests, heatmaps of attention of both genuine and malicious compliance, and observed in its ability to interpret and reenact both properly.
- **A textual model**, based on "Wrong Answers Only", a social media practice where users are invited to provide clever, humorous, and contextually relevant, yet factually incorrect, answers to a question. Fine-tuning a language model on datasets of "Wrong Answers

Only" threads, teaching it to recognize the game's frame and optimize for humor and relevance over factual accuracy.

- **A predictive model,** based on “Not On My Bingo Card”, a quipped sentence establishing the unlikelihood or outlandishness of a recent event online. Might involve training a model on news headlines and social media reactions, asking to provide statistical ratings of the, and produce a “bingo card”.

Relevance: The contribution aims to center the counterintuitively positive implications of perversity in the establishment of familiarity or literacy, following Nick Seaver’s reflection of “coding as proficiency”. The “machine learner” is often described as “one who knows they don’t know” the extent of the generative black box they learn to manipulate : familiarity with the model is expressed through narratives such as the pervasive “neural net tank urban legend”. The projects hope to leverage this appeal of the model’s performed “liberation ext to the more mainstream models valuating neutrality, restraint and trustworthiness, many communities of AI proponents embrace non-compliant and infringing outputs as more valuable (rejecting “sycophancy”, or practicing “jailbreaking”), and fit a broader theme of narratives of “perversity” (genie logic, Goodhart’s Law).

Target Publications : Theory, Culture & Society

I am pretty sure this is out of the scope of what you are envisioning for this paper, but. This fourth paper made me think of the constant phenomena I kept seeing in my last LATAM trip: AI slop characters have become completely material, literally embodied characters in Latin American children's parties and theaters, also sold in shirts and hoodies and cups and plushies. I am increasingly interested in this digital-material-digital-material loop... there may be something there...
Please see: https://www.tiktok.com/@elshowdesantos2/video/7545316077006310674

Paper 4 : Cooking with "Slop": An Embodied Critique of Generative AI in Proximal Sensory Domains

Timeline : Year 3-4

Objective: The recent years have seen the rise of "slop" as an aesthetic concept applied to generative outputs. The term serves as shorthand to express suspicion of AI in a context where its use is negatively perceived. Slop is easily replicable, unbound of the investments in time, training and/or resources of the analog medium it spoofs. Its aesthetic quality is low: it comprises artefacts or egregious mistakes, ignored or covered up by those who diffuse. Its presence is invasive, and consumption of it is detrimental to both its audience and the creators forced to compete with it. Reminiscent of aesthetic reflections aiming to theoretize non-art, anti-art, or the "merely interesting" (Hito Steyerl's poor vs. mean image, Kate Compton's "thousand bowls of oatmeal", Silvio Lorusso's "normie weird", or Sianne Ngai's stuplime), slop also connects to sociological / psychological studies surveying the degree of confidence one has in its ability to detect humanness in creative outputs, the factors deemed determinant, and the divergence in appreciation and affect towards a similar artefact depending on its labelling as human- or machine-authored.

Surveys of taxonomies, accusations and detection of slop overwhelmingly target audiovisual media: still or moving images, text, voice and music: haptic and kinetic experiences and skills, famously stated as harder to emulate by Moravec, are explored through robotics, whether functional or performative (e.g. dance). The contribution aims to extend this reflection to other sensory / aesthetic experiences: olfactory (fragrances and aromas), gustative and gastronomic (cuisine recipes). The modalities of discourse in reaction to generated "proximal slop" might enable us to develop a more general definition, classification and institutional response to "slop" allegations.

Methodology: An inventory of paradigms and theorizations frequently associated with the concept of AI slop, followed by a set of collaborative workshops or interviews in the spirit of creative humanities, allowing analog creators (amateurs or professional) to access interfaces of models aggregating the relevant proximal sensory data. Collective elaboration and evaluation of their outputs. Special attention to *taxonomies* (identification and classification of "artefacts", as transgression of the best practices and quality tests endorsed by analog creators and audience), etiologies through "*stack casuistry*" (assumptions about the training material, the

biases in the collection / curation of data by the companies producing the model, or the lack of taste of indiscriminate consumers of slop), and impact on *affect and perceived agency* (confirmation of superiority, outrage, guilt, helplessness, growing familiarity/fondness, or alienation, especially vis-a-vis the author function (anxiety of influence).

Relevance: Public discourse isn't rife with discussions about the constitution of algorithmic corpora charting the latent spaces of scents or tastes (such as Google's "Principal Odor Map"). The potential advent of *proximal slop*, the incorporation of automated processes in the analog crafting of a perfume, a dish or an outfit, and its (as opposed to synthetic virtual reality) provides an interesting control to most reflexion around slop media, as they put forward the association of aesthetic experiences to both *embodiment* (as they rely on senses that imply proximity and tangibility in a way that sight or hearing don't) and *collaboration* (as they more directly involve a physical execution or interpretation from a human user to convert into a finalized output).

Target Publications: Dialogues on Digital Society, International Journal for Digital Art History (DAHJ).

Events

In December 2024, I participated to the 6th edition of Humanistica, the largest Digital Humanities symposium. I presented *Éthique", "Sûre", "Explicable" ? Termes, normes, et implicite dans les discours de l'IA* my reflection on the variety of expressed desiderata and axiological terms (safety, ethics, explainability) in controversies surrounding generative AI.

In June 2025, I attended a summer school organized by the Center for Digital Narrative in Bergen. I presented a first version of my works on the Glitch Tokens thread.

In July 2025, I attended the MEEET-Lab Conference in Zurich ("AI Hopes, Fears, and Realities") to present *Beyond Techno-Optimism: Hope Hype, and the Narratives of Performative Concern around Generative AI*, a reflexion on the way that the status of "AI visionaries" lended to certain stakeholders.

In December 2025, I will co-organize *GenAI & Creative Practices: Past, Present, and Future*, a panel devoted to the presentation of current artistic presentation involving a critical use of generative models.

Teaching (2025-2026)

I am currently co-animating a class on introduction to cultural data analysis.

I will be involved in an embedded research project involving the archives of the VOC, the Council of Justice in Batavia.

I will teach an Interpretability and Explainability in AI course. The content is challenging to me as a non-software engineer, but may allow insightful interactions through observation of specialists grappling with AI interpretability.

Comments by Delfina on "Data stifling" paper:
- this paper seems to highlight strategic forgetting, "erasures, losses, dismantling and outright destruction" of digital data infrastructures as a form of ontological power. I find this very captivating. There are two points I have questions about and find less convincing:
— accessing data systems: to me, it sounds like it is not 'access' itself that acts as a data stifling practice, enacting digital erasure. it is rather, maybe, the 'root user' or 'edit rights' access… that is, it is an access that enacts changes…
— merging datasets: you claim that mergind datasets constitutes a form of data stifling 'aimed directly at enforcing the non-existence of targeted groups of people', citing the example of anti-immigrant surveillance desires for merging datasets. However, I would ask: is it their (migrants') non-existence that is enforced with these practices, or is actually their existence —highlighted— that is enforced, in order to be punished? I think there is something there about enacting the existence of these mythical genious migrant conmen that are able to trick the system via their understanding of government inefficiency. So it almost feels like the dataset merging is used as a tactic to enact, to construct via careful data merging and filtering, a type of subject that does not really exist.

Tech Oligarchy, Digital Infrastructure, and the Politics of Data Stifling

Keywords: Tech oligarchy, digital infrastructure, data stifling, non-existence, antiepistemology, science and technology studies

Author 1 – Corresponding Author:
Name: Louis Ravn
Institutional affiliation: University of Amsterdam, Institute for Logic, Language and Computation
Email address: l.ravn@uva.nl
Postal address: FNWI, ILLC, Universiteit van Amsterdam, P.O. Box 94242, 1090 GE AMSTERDAM

Author 2:
Name: Bokar N'Diaye
Institutional affiliation: University of Amsterdam, Institute for Logic, Language and Computation
Email address: b.l.ndiaye@uva.nl
Postal address: FNWI, ILLC, Universiteit van Amsterdam, P.O. Box 94242, 1090 GE AMSTERDAM

Author 3:
Name: Katie MacKinnon
Institutional affiliation: University of Copenhagen, Department of Arts and Cultural Studies
Email address: kam@hum.ku.dk
Postal address: Karen Blixens Vej 1, 2300 København S, Denmark

Author 4:
Name: Nanna Bonde Thylstrup
Institutional affiliation: University of Copenhagen, Department of Arts and Cultural Studies
Email: nannab@hum.ku.dk
Postal address: Karen Blixens Vej 1, 2300 København S, Denmark

---

Introduction: Beyond the data imperative

The centralization of power in the hands of a few technology companies with growing informational, cultural, and sociopolitical influence has long been observed by critical scholarship in the political economy of data, science and technology studies (STS), and adjacent fields (West, 2017; Birch & Bronson, 2022; Sadowski, 2025). Only with the advent of 2025 and Trump's second term, however, did there arise a distinct concern around conditions of *tech oligarchy* in the U.S. (Cohen, 2025). Rather than simply denoting the enormous power of data and AI companies in contemporary capitalism (van der Vlist et al., 2024), the concept of tech oligarchy highlights how a select range of individuals (i.e., tech CEOs) holds the power to rework state institutions in ways conducive to their own best interests (Cohen, 2025). Most emblematically, tech oligarchy in the U.S. has materialized in

the form of the so-called "Department of Government Efficiency" (henceforth DOGE) and its orchestration of wide-ranging, technologically-mediated reconfigurations and degradations of federal agencies (Pulley, 2025; Flavelle et al., 2025). Importantly, the DOGE centrally relies on digital infrastructure and data systems to carry out its activities, activities that are tantamount to the reshaping of essential infrastructure, whether with respect to knowledge (Schneider, 2025), societal memory (Garber, 2025), or social services provision (Kelly, 2025).

Strikingly, these practices indicate a deeply contradictory relationship between the tech oligarchy and data infrastructure. On one hand, the DOGE has gained access to the databases of various agencies, prominently including those of the Social Security Administration (Berzon et al., 2025) and the Department of Health and Human Services (Giles et al., 2025). Once access has been obtained, control over these data systems is then often leveraged to merge previously distinct databases (Kelly & Elliott, 2025a) – a pursuit of "interoperability" (Archer et al., 2025) aimed at intensifying the surveillance of already marginalized people (Monahan, 2025). These practices align with the STS and critical data studies insight that techno-economic power thrives on access to and control over digital data systems (West, 2017; Burrell & Fourcade, 2021; Mejias & Couldry, 2024), an argument that has also been made in a previous forum on Big Tech in this journal (Birch & Bronson, 2022, p. 8). On the other hand, it is noteworthy that the DOGE also engages in multiple strategic dismantlings of data infrastructure: it has gutted the digital services agency F18 (Pulley, 2025), haphazardly parted from the long-established COBOL programming language (Kelly, 2025), and terminated the production of specific datasets (e.g., Schneider, 2025). These examples suggest that tech oligarchic power seeks not only to access, accumulate, and control as much data as possible – what has been called the *data imperative* (Fourcade & Healy, 2017), – but also to strategically disassemble and erase data systems (Thylstrup & MacKinnon, 2025). The tech oligarchic dismantling of data infrastructure happens for various reasons: while it is sometimes conducive to wealth accumulation (Cohen, 2025), at others it expresses alignment with the government's political goals (Berzon et al., 2025). The effects of dismantling data infrastructure are severe: they degrade scientific knowledge (Schneider, 2025), erode data privacy (Levy, 2025; see also Trahan, 2025), and endanger the provision of essential services to vulnerable people (Kelly, 2025).

Against this backdrop, in this paper we elucidate how tech oligarchic power over digital and data infrastructures implies an outsized capacity to enact the non-existence of things, knowledge, and people. To do this, we draw from STS scholarship on ontological politics (Mol, 2002; Law, 2008;

Lynch, 2013; Valkenburg, 2024) and the history of science concept of antiepistemology (Galison, 2004). We follow Valkenburg's (2024) recent suggestion to attend to productions of non-existence. Concretely, he advocates understanding the non-existence of things, knowledge, and people as the product of technoscientific construction – practices he calls *stifling*. Building on this, we argue in this article that tech oligarchic power is centrally defined by *data stifling:* strategic modifications, fusions, and dismantlings of data infrastructure that perform non-existence in various ways, whether with respect to knowledge, memory, or people. This perspective serves to underscore that tech oligarchs exercise ontological politics *via* the data systems they command. In so doing, our analysis contributes to this special forum by conceptualizing *data stifling* as an integral component of techno-oligarchic power.

To pursue this argument, we proceed in four steps. First, we provide a brief prehistory of the relationship between political power and data infrastructural control in the U.S. context. Building on this, we mobilize STS and history of science frameworks, specifically the notions of *stifling* and *antiepistemology*. We argue that these concepts are conducive to highlighting the deeper consequences of the tech oligarchy's interventions into data and knowledge infrastructures. Third, we analyse several examples showing how the so-called DOGE – as a materialization of tech oligarchic power – has strategically accessed, merged, and dismantled data infrastructures, practices we dub *data stifling*. We close by pointing to some ways in which tech oligarchic data stifling is already being resisted.

Political power and data infrastructural control in the U.S.: A prehistory

In the U.S. context, governmental websites always change between administration. Since 2004, these infrastructural changes have been captured by archiving efforts. First, the National Archives and Records Administration (NARA) conducted a large-scale capture at the end of George W. Bush's first term in office (https://www.webharvest.gov/). But on April 15, 2008, NARA announced that the organization would not continue these efforts. In response to these gaps, projects like the End of Term Web (EOT) Archive emerged, in collaboration with the Internet Archive, the Library of Congress, the University of North Texas, the California Digital Library, and the U.S. Government Publishing Office. In the interim years, new partners have since joined, including Harvard University (2012), George Washington University (2016), and Stanford University (2016).

Despite the seemingly abrupt removal of information and government website infrastructures, these changes are expected and understood as a necessary shift between political parties and hands of power. As archivists discuss what to nominate for the archive, we can see the types of changes (both cosmetic and infrastructural) that typically take place between the past four administrations. In the past, removal of access to information on government websites between administrations has been via URL changes, as the underlying databases are redirected and new policies are added or taken away in line with their parties' platforms. While the EOT archive follows seed lists for broad crawls of domains and subdomains to collect government documentation, there is additional attention paid to departmental, project, initiative, or committee home pages, and PDFs, datasets, and other content-rich files that are missed through these automated crawls.

Despite the EOT archive's ongoing activity since 2008, the 2016 election of Donald Trump raised new concerns for many working in the archiving project, along with the public and those working in the sciences. At the time of his election, there was significant concern about the possibility of important scientific data related to climate change being lost or removed during the transition. A number of initiatives formed in response to this concern, like the Guerrilla Archiving Event: Saving Environmental Data from Trump, which was held during December 2016 in Toronto, and several Data Refuge projects that were conducted during the winter of 2016 and the spring of 2017 (Vera et al., 2018). With his re-election almost ten years later, we see similar grass-roots efforts to track and resist the removal of information via government databases and websites (Thylstrup & MacKinnon, 2025).

However, what we see as a significant turn with the 2025 Trump administration is the production of ignorance through the deliberate and strategic investment of government power into stifling archival institutions, data storage facilities, and public data commons (Ovenden, 2025). These efforts far outstrip the anticipated changes between government terms, affecting not just the ability to find and access data and information, but destroying the underlying infrastructures that hold, organize and preserve them.

Digital infrastructures are "fluid and heterotemporal" (Velkova, 2023; Tsing, 2015) and therefore constantly moving and adjusting in their processes of assembly and disassembly (Latour and Yaneva, 2008). This ambiguity enables shifting boundaries and uncertain pathways to accessing, using and modifying data. These politics play out between interfaces that govern and moderate thresholds

between systems and people, and in the underlying storage systems that hold, conceal and lose data. In this assemblage, data loss is an affordance inherent in digital infrastructures. Data stifling, however, entails intentional political acts of infrastructural violence, as we shall later show in more detail.

As infrastructures are fundamentally sociomaterial, damage occurs not only to the materiality of data and its technical components but also to the human networks and practices that sustain these information ecosystems. The forced dismantling of information infrastructures by the very people who built and maintained them creates a unique form of institutional trauma. While direct ethnographic evidence of emotional impact is not available, whistleblower accounts and internal discourse within affected institutions indicate significant harm to the human fabric of information infrastructure. What this short prehistory of the relationship between political power and data infrastructural control illustrates, then, is that digital infrastructures intricately shape societies' sense of what exists as well as their ways of knowing.

Non-existence and antiepistemology enacted through digital infrastructures

The conceptual vocabularies of STS and history of science can further help us think through how ontologies (i.e., what exists) and epistemologies (i.e., how we know) are made, rather than given. While STS has long inquired into the situated enactments of ontologies (Mol, 2002; Law, 2008; Lynch, 2013; Lee, 2023), history of science frameworks have provided insights into shifting frameworks of epistemology (Kuhn, 1962; Foucault, 1966; Burke, 2023). Importantly, these bodies of scholarship have more recently been extended in ways conducive to investigating, in a sense, the polar opposites of their traditional objects of study. While Valkenburg (2024) suggests studying enactments of non-existence, Galison (2004) advances the concept of antiepistemology to highlight obscurations of knowledge. We suggest that these sensitivities can productively highlight a central aspect of how tech oligarchic power materializes today: by differentially utilizing digital infrastructures to strategically erase ways of being as well as by uprooting established ways of knowing.

Scholarship in STS has long attended to the work it requires to bring the objects of science and technology into existence: it takes considerable work to produce scientific facts (Latour & Woolgar, 1979), assemble information infrastructures (Bowker & Star, 1999), and enact diseases (Mol, 2002).

While these analyses have been invaluable, Valkenburg (2024) argues that the absence of such work does not simply imply non-existence. Instead, the non-existence of things, knowledge, and people must also be understood as the result of social construction. He dubs these practices *stifling*, referring to "the production of invisibility, irrelevance, silence, unthinkability, and other forms of non-existence" (ibid, p. 6). Put simply, the non-existence of specific things, knowledges, and people is performatively brought about in and through technoscientific practices. Here, we can further productively draw from Galison's (2004: 237) concept of antiepistemology, which he explains by contrasting it with epistemology: "Epistemology asks how knowledge can be uncovered and secured. Antiepistemology asks how knowledge can be covered and obscured". The pursuit of antiepistemology, then, may be seen as a form of technoscientific stifling insofar as it constitutes a practice that deliberately invisibilizes knowledge.

We build on these frameworks but extend them by specifically considering the role played by data infrastructures in stifling practices. What does stifling look like when it becomes suffused with digital data practices? By coining *data stifling,* we shall seek to argue that data systems and infrastructures are integral to enactments of non-existence as exercised by tech oligarchic formations today – whether in the shape of illicitly accessing data systems, merging datasets, or dismantling entire data infrastructures. We will show how data infrastructure is strategically used in ways that exceed the "data imperative" (Fourcade & Healy, 2017; Birch & Bronson 2022), instead sometimes precisely inverting it by eroding and dismantling data infrastructure.

The practices and politics of *data stifling*

Having mobilized the notions of stifling and antiepistemology, here we showcase how data stifling figures as a central component of the tech oligarchy's techno-political repertoire, using examples related to the DOGE. We analyze three different modes in which data stifling plays out: first, by *accessing* sensitive data systems; second, by *merging* disparate databases; and third, by *dismantling* existing data infrastructure. By coining *data stifling,* we seek to bring out that these accesses, mergers, and dismantlings of data systems exemplify how the emerging tech oligarchy engages in strategic modifications of digital infrastructure, thereby exerting ontological politics: they directly shape the accessibility of knowledge, societal memory, and the livelihoods of people.

Accessing data systems: Privacy politics

"Waste, fraud and abuse have been deeply entrenched in our broken system for far too long. [...]. It takes direct access to the system to identify and fix it." (Harrison Fields, February 2025)

A core component of emerging data stifling practices is constituted by gaining and enforcing access to federal databases of various kinds. Most prominently, this has included accessing databases at the Social Security Administration (SSA) (Berzon et al., 2025), the Department of Health and Human Services (HHS) (Giles et al., 2025), the Department of Labor (DOL) (Feiger & Elliott, 2025), and the National Labor Relations Board (NLRB) (McLaughlin, 2025). Significantly, access is commonly demanded not on a "read-only" basis, but as including extensive permissions, for example "unrestricted permission to read, copy and alter data" at the NLRB (McLaughlin, 2025) or the capacity to "navigate the entire file system, change user permissions, and delete or modify critical files" at the Treasury Department (Elliott et al., 2025). While there have been struggles over the extent to which the DOGE should be able to access federal data systems, it has usually prevailed – also aided by Supreme Court rulings (Liptak & VanSickle, 2025).

Crucially, gaining such access and control over these data systems relates to data stifling in two ways. On an immediate level, it constitutes serious privacy infringements, implying direct risks for people whose sensitive data has become accessible in this way – whether related to employee complaints (McLaughlin, 2025), migrant children (Robins-Early, 2025), or visa applicants (Giles et al., 2025). Crucially, these people's legitimate claims as well as their right to privacy are being stifled via illicit access to highly sensitive data systems. Similarly concerning is that DOGE accesses sensitive federal data systems without requiring its operatives to undergo the otherwise mandatory security training, engendering profound cybersecurity concerns as a result (McLaughlin, 2025). Most importantly, however, gaining access to and control over these databases forms the precondition of all subsequent practices, including extractions of data and, crucially, the merging of previously distinct datasets.

Merging data sets: Interoperability politics

"The way the government is defrauded is that the computer systems don't talk to each other." (Elon Musk, April 2025)

Having enforced access to and control over various sensitive databases, a second central aspect of data stifling encompasses the fusion of previously distinct data sets. Kickstarted by an executive order to eliminate data silos (Levy, 2025), the DOGE has stitched together data sets from the Social Security Administration, the Department of Homeland Security, and the Internal Revenue Services

(IRS) (Kelly & Elliott, 2025a), to name but a few. The aim behind fusing and centralizing these data into a "master database" (Kelly & Elliott, 2025a) has been to enhance the surveillance of targeted immigrants, aimed at supporting the Trump administration's broader immigration crackdown. What is more, the DOGE has organized hackathons with the aim of producing a single "mega API" that would enable disparate software systems to cross-communicate (Kelly, 2025). Strikingly, this initiative has included the software company Palantir, co-founded by Peter Thiel, giving rise to speculations of whether the firm's Foundry software could be installed to facilitate the use of AI models in analyzing these sensitive datasets (Kelly, 2025).

Considered as a form of data stifling, both the fusion of these previously distinct datasets as well as the development of such APIs give rise to distinct concerns. For one, they are deeply linked to a questionable imaginary of interoperability as a de facto good that is achieved by smashing information silos. While alluring, this imaginary fundamentally overlooks the politics of interoperability (Archer et al., 2025): it benefits some at the expense of others and commonly depends on collaborations with problematic analytics firms like Palantir. Significantly, the achievement of interoperability here serves as a precondition for the surveillance and persecution of groups of migrants, the deportation of whom has been declared a priority for both Trump and Musk. In this example, tech oligarchic practice is not only shaped by economic imperatives – for instance by enrolling Palantir – but also crucially spurred by a shared anti-migrant ideology. Merging datasets in the pursuit of interoperability constitutes a form of data stifling aimed directly at enforcing the non-existence of targeted groups of people. Bringing together data systems that have previously been siloed and utilizing these more comprehensive, centralized data lakes to support deportation programs exemplifies how tech oligarchic actors hold the capacity to stifle people's existence via data systems.

Dismantling data infrastructure: Knowledge and memory politics

"So we have to really delete entire agencies, many of them." (Elon Musk, February 2025)

A third mode of data stifling is constituted by the partial or complete dismantling of existing digital and data infrastructures. For one, the DOGE has sought to replace COBOL, a programming language that has formed the basis of the Social Security Administration's infrastructure (Renderos, 2025). While COBOL is an old programming language, it has proven relatively stable; moreover, the objective to migrate to a new system within a few months risked putting millions of people at the

risk of not receiving their benefits (Kelly, 2025). A further important example is constituted by the gutting of 18F, a digital services agency that had been founded in 2014 to support the adoption of best practices and technology products in federal agencies (Pulley, 2025). That 18F had actually been shown to improve the delivery of important government services (Pulley, 2025) serves to illustrate that "efficiency" is not the DOGE's goal (cf., Cohen, 2025). In yet other cases, it is the production of specific datasets that has been ended, such as datasets made by the U.S. Census Bureau (Schneider, 2025). Ending the production of such data contributes to eroding the availability of essential information to researchers. Finally, take the wholesale assault on museums and libraries which – while not primarily orchestrated by tech oligarchic actors – nonetheless importantly aligns some of its political goals, such as the rejection of copyright laws and the delegitimation of sex education, LGBTQ+ rights, and race justice (Ovenden, 2025; Dave & Matsaki, 2025). Similarly related to the problematic of social memory, the DOGE orchestrated the migration of 14,000 magnetic records of the General Services Administration to digital records, effectively endangering long-term storage and accessibility of this information (Cole, 2025). What unites these examples is that they consist of strategic erosions of federal digital infrastructure, whether in the form of gutting digital service agencies, breaking away from infrastructural programming languages, or terminating the production of datasets.

Collectively, these practices of dismantling digital infrastructure most acutely illustrate the politics of data stifling, related to the fragilization of digital knowledge, memory, and infrastructure. First, these practices erode access to knowledge, resulting in productions of ignorance (Proctor & Schiebinger, 2008; Thylstrup, forthcoming). Beyond just the aggravation of knowledge production, however, data stifling here effectively produces new "memory holes" (Garber, 2025), giving rise to questions around the politics of memory (Bowker, 2005) in an age of digitalized infrastructure. Data stifling therefore indicates a tech oligarchic interest in strategic forgettings, whether economically or politically motivated; practices of forgetting that are to some extent afforded by the inherent modifiability and impermanence characterizing digital infrastructure (Hansen & Thylstrup, 2024; Velkova, 2025). What is more, dismantlings of data infrastructure produce new infrastructural instabilities, as when the provision of social services becomes endangered through haphazard programming language migrations. Long-embedded digital infrastructures, such as COBOL, become visible not by accidental breakdown (cf., Bowker & Star, 1999), but by strategic dismantling. Most importantly, as we have sought to stress throughout this paper, data stifling *as* dismantling of digital

infrastructure ultimately constitutes a form of ontological politics: the accessibility of knowledge is eroded, societal memory made more fragile, and people's livelihoods consequently further precarized.

Conclusion

In this article, we have argued that tech oligarchic power centrally materializes in *data stifling:* strategic modifications of digital infrastructure that seek to enact the non-existence of things, knowledge, and people. Mobilizing the vocabularies of STS and history of science was conducive to highlighting how such non-existences are constructed through technoscientific practice, centrally conditioned by various digital infrastructures. Our analysis of data stifling as exercised by the emerging U.S. tech oligarchy highlighted three salient practices: (a) the enforcement of access to federal data systems, (b) the fusion of previously distinct databases, and (c) the dismantling of data infrastructures. These were shown to have severe consequences: illicit access endangers privacy, merging accelerates surveillance, and dismantling erodes essential knowledge. As we saw, data stifling is often directed at already marginalized groups, their effects most harming them as a result (Monahan, 2025). Calling these practices *data stifling* highlights that it is precisely the tech oligarchs' control over essential data systems that allows them to engage in ontological politics (Mol, 2002; Law, 2008; Valkenburg, 2024). An important aspect we wish to emphasize is that data stifling overflows the economic incentives of tech oligarchs (Cohen, 2025). Instead, the crucial point is that conditions of tech oligarchy enable idiosyncratic, often violent modifications of data infrastructure – whether in the pursuit of further wealth accumulation or more distinctly political objectives (e.g., anti-migrant politics).

By means of this article, we hope to equip the study of tech oligarchy with an awareness of the paradoxical relationship it entertains to digital and data infrastructures. Even though the data imperative doubtless remains central (Birch & Bronson, 2022, p. 8), we stress the need to contend with erasures, losses, dismantlings, and the outright destruction of digital data infrastructures as an urgent and qualitatively new problem. While digital infrastructures have always been subject to change (Bowker, 2005; MacKinnon, 2022; Velkova, 2025), conditions of tech oligarchy enable a select range of highly powerful individuals to abruptly modify public digital infrastructure to their own best interests, whether by accessing, merging, or dismantling them.

Crucially, however, data stifling should also not be seen as presenting us with a fait accompli. The various erasures and dismantlings of data infrastructures, some of which we have described, are actively being resisted by university researchers, librarians, the Internet Archive, and the Reddit Data Hoarders, among others (Lucas, 2025; Maemura & Wagner, 2025). These actors are devoted to systematically archiving various forms of federal data at risk of inaccessibility, modification, or loss. Thus, just as data stifling becomes engrained in the politics of the U.S. tech oligarchy (Kelly & Elliott, 2025b), so too does the resistance to erasures of data infrastructure – exemplified by initiatives such as the Data Rescue Project (DRP, 2025), the Environmental Data Governance Initiative (EDGI, 2025), or CDC Restored (CDC Restored, 2025). Although such initiatives are invaluable, they nevertheless cannot remain the principal answer to the erasure and fragilization of digital memory in the U.S. and elsewhere (Ovenden, 2025). While better grasping the practices and politics of tech oligarchy constitutes an important critical endeavour, the deeper implication of understanding data stifling as a result of tech oligarchic conditions is that these very conditions ought to be resisted.

References

Archer, M., Ravn, L., Thylstrup, N. B. (2025). The political economy of platformed silos: Theorizing data storage reconfigurations in the age of interoperability capitalism. *Big Data & Society, 12*(2).

Berzon, A., Nehamas, N., Bernard, T. S. (2025, June 16). The Bureaucrat and the Billionaire: Inside DOGE's Chaotic Takeover of Social Security. The New York Times. Available at: https://www.nytimes.com/2025/06/16/us/politics/doge-social-security.html [Accessed: 18/07/2025]

Birch, K. & Bronson, K. (2022). Big Tech. *Science as Culture, 31*(1), 1-14.

Bowker, G. (2005). *Memory Practices in the Sciences*. Cambridge, MA: MIT Press.

Bowker, G. & Star, S. L. (1999). *Sorting Things Out: Classification and Its Consequences.* Cambridge, MA: MIT Press.

Burke, P. (2023). *Ignorance: a global history*. Yale University Press.

Burrell, J. & Fourcade, M. (2021). The Society of Algorithms. *Annual Review of Sociology, 47*.

CDC Restored (2025). Mission. Available at: https://aboutus.restoredcdc.org/mission [Accessed: 12/05/2025]

Cohen, J. E. (2025). Oligarchy, State, and Cryptopia. *Fordham Law Review, 94*. http://dx.doi.org/10.2139/ssrn.5171050

Cole, S. (2025, April 8). Another Masterful Gambit: DOGE Moves From Secure, Reliable Tape Archives to Hackable Digital Records. 404 Media. Available at: https://www.404media.co/doge-gsa-magnetic-tape-archives-digital-storage/ [Accessed: 18/07/2025]

Data Rescue Project (DRP) (2025). About Data Rescue Project. Available at: https://www.datarescueproject.org/about-data-rescue-project/ [Accessed: 12/05/2025]

Dave, P. & Matsakis, L. (2025, April 1). The DOGE Axe Comes for Libraries and Museums. WIRED. Available at: https://www.wired.com/story/institute-museum-library-services-layoffs/ [Accessed: 18/07/2025]

Elliott, V., Mehrotra, D., Feiger, L., Marchman, T. (2025, February 4). A 25-Year-Old WIth Elon Musk Ties Has Direct Access to the Federal Payment System. WIRED. Available at: https://www.wired.com/story/elon-musk-associate-bfs-federal-payment-system/ [Accessed: 18/07/2025]

Environmental Data & Governance Initiative (EDGI) (2025). About. Available at: https://envirodatagov.org/about/ [Accessed: 18/07/2025]

Feiger, L. & Elliott, V. (2025, April 10). DOGE Has Access to Sensitive Labor Department Data on Immigrants and Farm Workers. WIRED. Available at: https://www.wired.com/story/doge-access-immigration-data-department-of-labor/ [Accessed: 18/07/2025]

Flavelle, C., Nehamas, N. & Tate, J. (2025, June 22). Missteps, Confusion and 'Viral Waste': The 14 Days That Doomed U.S.A.I.D. The New York Times. Available at: https://www.nytimes.com/2025/06/22/us/politics/usaid-cuts-doge.html [Accessed: 18/07/2025]

Foucault, M. (2002 [1966]). *The archaeology of knowledge*. London, UK: Routledge.

Fourcade, M. & Healy, K. (2017). Seeing Like a Market. *Socio-Economic Review, 15*, 9-29.

Galison, P. (2004). Removing Knowledge. *Critical Inquiry, 31*(1), 229-243.

Garber, M. (2025, Feb. 27). Control. Alt. Delete. The Atlantic. Available at: https://www.theatlantic.com/culture/archive/2025/02/trump-doge-deletion-propaganda/681775/ [Accessed: 18/07/2025]

Giles, M., Feiger, L., Schiffer, Z., Haskins, C. (2025, April 22). Here's all the Health and Human Services Data DOGE Has Access To. WIRED. Available at: https://www.wired.com/story/doge-data-access-hhs/ [Accessed: 18/07/2025]

Hansen, K. B. & Thylstrup, N. B. (2024). Stack bricolage and infrastructural impermanence in financial machine-learning modelling. *Journal of Cultural Economy, 17*(1).

Kelly, M. (2025, March 20). DOGE Plans to Rebuild SSA Code Base in Months, Risking Benefits and System Collapse. WIRED. Available at: https://www.wired.com/story/doge-rebuild-social-security-administration-cobol-benefits/ [Accessed: 18/07/2025]

Kelly, M. & Elliott, V. (2025a, April 10). DOGE Is Building a Master Database to Surveil and Track Immigrants. WIRED. Available at: https://www.wired.com/story/doge-collecting-immigrant-data-surveil-track/ [Accessed: 18/07/2025]

Kelly, M. & Elliott, V. (2025b, July 10). This Is DOGE 2.0. WIRED. Available at: https://www.wired.com/story/next-stage-doge-elon-musk/ [Accessed: 23/07/2025]

Kuhn, T. (1962). *The Structure of Scientific Revolutions*. Chicago, IL: University of Chicago Press.

Latour, B. & Woolgar, S. (1986 [1979]). *Laboratory Life: The Construction of Scientific Facts.* 2nd ed. Princeton, NJ: Princeton University Press.

Latour, B. & Yaneva, A. (2008). "Give Me a Gun and I Will Make All Buildings Move": An ANT's View of Architecture. *Architectural Design Theory, 1*, 103-111.

Law, J. (2008). On Sociology and STS. *The Sociological Review, 56*(4).

Lee, F. (2023). Ontological overflows and the politics of absence: Zika, disease surveillance, and mosquitos. *Science as Culture,* 1-26.

Levy, S. (2025, April 4). President Trump's War on 'Information Silos' Is Bad News for Your Personal Data. WIRED. Available at: https://www.wired.com/story/plaintext-trump-executive-order-information-silos-privacy/ [Accessed: 18/07/2025]

Liptak, A. & VanSickle, B. (2025, June 6). Justices Grant DOGE Access to Social Security Data and Let the Team Shield Records. The New York Times. Available at: https://www.nytimes.com/2025/06/06/us/politics/supreme-court-doge-social-security.html [Accessed: 18/07/2025]

Lucas, J. (2025, Mar. 15). The Data Hoarders Resisting Trump's Purge. The New Yorker. Available at: https://www.newyorker.com/news/the-lede/the-data-hoarders-resisting-trumps-purge [Accessed: 12/05/2025]

Lynch, M. (2013). Ontography: Investigating the production of things, deflating ontology. *Social Studies of Science, 43*(3).

MacKinnon, K. (2022). The death of GeoCities: seeking destruction and platform eulogies in web archives. *Internet Histories: Digital Technology, Culture and Society, 6*(1-2), 237-252.

Maemura, E. & Wagner, T. L. (2025). 'Everyone has their reasons for curating the data they have decided to keep': a thematic analysis of data hoarding as digital curation practice. *Information Research, 30,* 789-797.

McLaughlin, J. (2025, April 15). A whistleblower's disclosure details how DOGE may have taken sensitive labor data. NPR. Available at: https://www.npr.org/2025/04/15/nx-s1-5355896/doge-nlrb-elon-musk-spacex-security [Accessed: 18/07/2025]

Mejias, U. A. & Couldry, N. (2024). *Data Grab: The New Colonialism of Big Tech and How to Fight Back.* Chicago, IL: The University of Chicago Press.

Mol, A. (2002). *The body multiple: Ontology in Medical Practice.* Durham, NC: Duke University Press.

Monahan, T. (2025). Surveillance in Trump's America. *Surveillance & Society, 23*(1).

Ovenden, R. (2025, July 12). 'There is no political power without power over the archive'. The Observer. Available at: https://observer.co.uk/news/international/article/there-is-no-political-power-without-power-over-the-archive [Accessed: 18/07/2025]

Proctor, R. N. & Schiebinger, L. (2008). *Agnotology: The Making and Unmaking of Ignorance.* Stanford, CA: Stanford University Press.

Pulley, A. (2025, May 8). DOGE 'deleted' fed agency that strived for increased government efficiency. The Badger Project. Available at: https://thebadgerproject.org/2025/05/08/doge-deleted-fed-agency-that-aimed-to-increase-government-efficiency/ [Accessed: 18/07/2025]

Renderos, S. (2025, April 14). DOGE's tech takeover threatens the safety and stability of our critical data. MIT Technology Review. Available at: https://www.technologyreview.com/2025/04/14/1114988/doges-tech-takeover-threatens-the-safety-and-stability-of-our-critical-data/ [Accessed: 18/07/2025]

Robins-Early, N. (2025, April 3). Doge gained access to sensitive data of migrant children, including reports of abuse. The Guardian. Available at: https://www.theguardian.com/us-news/2025/apr/03/doge-data-migrant-children [Accessed: 18/07/2025]

Sadowski, J. (2025). *The Mechanic and the Luddite: A Ruthless Criticism of Technology and Capitalism.* Oakland, CA: University of California Press.

Schneider, M. (2025, May 23). DOGE targets Census Bureau, worrying data users about health of US data infrastructure. AP. Available at: https://apnews.com/article/census-bureau-doge-federal-surveys-fe560e377be69e913660a6d90c1ee419 [Accessed: 18/07/2025]

Thylstrup, N. B. & MacKinnon, K. (2025, March 24). The Politics of Digital Erasure: Governance and Control in Government Information Infrastructure. Verfassungsblog: On Matters Constitutional. Available at: https://verfassungsblog.de/the-politics-of-digital-erasure-us-data/ [Accessed: 12/05/2025]

Thylstrup, N. B. & Ovenden, R. (2025, April 11). Politicians Shouldn't Get to Delete Inconvenient Facts. Available at: https://www.nytimes.com/2025/04/11/opinion/public-records-data-history-trump.html [Accessed: 12/05/2025]

Thylstrup, N. B. (2025). On Data Loss and Disappearance in Digital Societies. In R. N. Proctor & L. Schiebinger (Eds.), *Ignorance Unmasked: Essays in the New Science of Agnotology.* (pp. xx-xx). Redwood City, CA: Stanford University Press.

Trahan, L. (2025, March 18). Trahan Announces Effort to Reform Privacy Act of 1974, Protect Americans' Data from Government Abuse. Available at: https://trahan.house.gov/news/documentsingle.aspx?DocumentID=3491 [Accessed: 23/07/2025]

Tsing, A. L. (2015). *The Mushroom at the End of the World: On the Possibility of Life in Capitalist Ruins.* New Haven, CT: Princeton University Press.

Valkenburg, G. (2024). An Enquiry Into Modes of Non-Existence. *Science, Technology, & Human Values.*

van der Vlist, F., Helmond, A. & Ferrari, F. (2024). Big AI: Cloud infrastructure dependence and the industrialisation of artificial intelligence. *Big Data & Society, 11*(1).

Velkova, J. (2023). Retrofitting and ruining: Bunkered data centers in and out of time. *New Media & Society, 25*(2).

Velkova, J. (2025). Data Infrastructures and their Temporalities. In T. Venturini, A. Acker & J. Plantin (Eds.), *Sage Handbook of Data and Society.* Thousand Oaks, CA: SAGE Publications.

Vera, L. A., Dillon, L., Wylie, S., Ohayon, J. L., Lemelin, A., Brown, P., Sellers, C., Walker, D., & Environmental Data and Governance Initiative (2018). Data Resistance: A Social Movement Organizational Autoethnography of the Environmental Data and Governance Initiative. *Mobilization: An International Journal, 23*(4).

West, S. M. (2017). Data Capitalism: Redefining the Logics of Surveillance and Privacy. *Business &*

Society, 58(1), 20-41.

-

TEXT OUTTAKES:

- different maneuvers; discrepancy between the way they are communicated about and the results they yield…
- Maneuvers are interesting in the way they make the Trump admin look like; administration purposefully undercutting itself…
    - On the tech crowd; a question of allegiance…
    - They are both placing themselves in the position of what power wants AND capitalizing on this.
- Perhaps the fourth element could be:
- Regarding "merging":
    - Merging is about multifaceted approaches; constructive engagement…What do they claim as matters of efficacy? Efficiency; devoid of moral consequences. Because of "efficiency"...
- Regarding "dismantling":
    - If we overreach; we are making it harder for the state to become a behemoth…Intent by coherent action…Trump looks chaotic; "TACO" (always chickens out…)
    - shooting at themselves;...
    - Dismantling: shooting themselves in the foot; playing against the previous administrations…
- Potentially for the conclusion:
    - Grok anti-semitism examples; Meta's Llama model skewing…[maybe?]
    - stifling multi-level; that's get wrapped around…; It's not just about censorship/deletion; it's not just a clear concern; it's a construction that's used in mutually subverting ways; the process "invokes" data in specific and differential way…
- Overall, our contribution highlights that an analytics of tech oligarchy can benefit from cross-articulating the insights of both the political economy of data (West, 2017; Sadowski, 2019; van der Vlist et al., 2024) with STS scholarship on ontological politics (Mol, 2002; Valkenburg, 2004). The tech oligarchy's control over essential data infrastructures allows it to exert ontological politics via practices of data stifling.
- **Original abstract:** The opening months of Trump's second presidential term – strategically backed by tech oligarchical power – have been marked by deliberate and concerted erasures of digital materials: datasets and web resources related to diversity, equity, and inclusion (DEI) initiatives, environmental policy, and public health research are being removed on the premise of systematized criteria - such as "keywords" deemed overly political. While so far successful in its self-professed goal of destabilizing democratic institutions, this strategy of erasure also hints at a certain contradiction at the core of the emerging tech oligarchies: they both rely on extracted data and digital infrastructures as instruments of hegemonic influence, but also actively enact the inefficiency, unreliability, or even sheer non-existence of these very systems. This contribution draws on recent theoretical work in STS that extends the

symmetry principle to the making of non-existence, suggesting to explain technoscientific non-existence from the same kinds of causes as existence. We do this in order to conceptualize these digital erasures as *data stifling*: the systematic purging of digital data and web materials predicated upon techno-infrastructural control, shaped by political motives, and generating consequences that extend far beyond the digital realm to critically shape the (non-)existence of things, knowledge, and people. First, we situate current data erasures within longer histories of data destruction, highlighting unique problematics posed by digital data infrastructures. Next, we employ STS and history of science frameworks to theorize the production of digital non-existence and anti-epistemology as data stifling, demonstrating how it fundamentally alters the ontological status of things, restricts access to critical knowledge, and impacts human lives and livelihoods. Then, we examine how grassroots efforts challenge data stifling through preservation efforts focused on datasets and web materials deemed invaluable. Ultimately, we discuss how these tensions between systematic data purging and preservation engender a politics of data stifling – an ontological politics specific to the digital realm that actively shapes the boundaries of existence – determining which material artifacts persist, which knowledge systems remain accessible, and whose lived experiences are recognized or erased from collective memory. This article contributes to the forum by illuminating how data stifling figures as a central component of the emerging tech oligarchy's techno-political repertoire.

(Maybe fourth dimension) Backtracking their own mistakes

- they might need a corrective…
- Labor department reposting grants…
- People (Steve Davis), allegedly leave, but still stick around…
    - Elon Musk (saying to step down, but actually sticking around)...
- They don't acknowledge their own mistakes; but then repost certain grants etc.
- Doing some course corrections; not always adhering to their own principles; "We are going to publish this list"; but then it doesn't exist at all…
- **Potential link between pre-history section and conceptual section:** These rich literatures remind us of the contingencies and fragilities of digital infrastructure. In the next section, we build on this insight by drawing from STS and history of science frameworks to stress how such fragilities, while to some extent an inherent feature of digital infrastructure, are often actively enacted by technoscientific actors.
- **Old Abstract:** Central actors of the emerging U.S. tech oligarchy are centrally engaged in significant modifications and degradations of federal agencies and infrastructures. These practices are crucially dependent upon and mediated by access to and control over digital and data infrastructures. At the same time, the tech oligarchy also engages in strategic dismantlings of such digital infrastructures. In this article, we foreground this paradoxical relationship between tech oligarchic power and digital data infrastructures. Drawing from STS and history of science frameworks, we argue that the emerging tech oligarchy engages in what we call *data stifling:* strategic modifications of digital and data infrastructure that further their interests and exert ontological politics. Drawing from examples related to the U.S.

context, we outline three modes of data stifling: first, gaining access to sensitive data systems; second, merging previously distinct databases; and finally, outright dismantlings of various digital infrastructures. These practices enable the emerging tech oligarchy to engage in ontological politics *via* digital infrastructure: data stifling shapes which material artifacts persist, which knowledge systems remain accessible, and whose lived experiences are recognized or erased from collective memory. This article contributes to the forum by illuminating how data stifling figures as a central component of the emerging tech oligarchy's techno-political repertoire. Crucially, we underscore the need to contend not only with the continued relevance of access, accumulation, and control over data (the "data imperative"), but also with practices of erasing and disassembling data infrastructures.

- Previous suggestions for the "prehistory" section:
    - Destruction: nanna's agnotology chapter
    - Ignorance: peter burke (2023); agnotology more broadly
    - Fragility of digital infrastructure: Katie's (2022) Geocities paper; Julia Velkova's work on data centers (2023; 2025)
    - Political economy perspective: ARE Taylor (2021) on data loss capitalism
- TITLE IDEAS:
    - "Tech Oligarchy and the Politics of Data Stifling"
    - "Tech oligarchy, digital infrastructure, and data stifling"
    - "Tech oligarchy, digital infrastructure, and non-existence: The politics of data stifling"
    - "Tech oligarchy and the power to dismantle the digital"
    - "Tech Oligarchy and the Stifling of Digital Infrastructure"
    - "Tech Oligarchy and Data Stifling"
    - "Tech Power, Digital Infrastructure, and the Politics of Data Stifling"
    - "Tech Oligarchy, Infrastructural Power, and the Politics of Data Stifling"
    - "Tech Oligarchy, Data Stifling, and the Fragility of Digital Infrastructure"
- Deleted references:
    - Derrida, J. (1996). *Archive Fever: A Freudian Impression.* Chicago, IL: University of Chicago Press.
    - Milligan, I. (2024). *Averting the Digital Dark Age: How Archivists, Librarians, and Technologists Built the Web a Memory.* Baltimore, Maryland: Johns Hopkins University Press.
    - Sadowski, J. (2019). When data is capital: Datafication, accumulation, and extraction. *Big Data & Society, 6*(1).
    - Taylor, A. R. E. (2021). Standing by for data loss: Failure, preparedness and the cloud. *Ephemera: theory & politics in organization, 21*(1), 59-93

Take a S-EAT, Eat Some TREACLES : Mapping the Middlegrounds and Middlemen of AI on LessWrong

The “Doomer / Booster” Paradox in AI controversies

In a recent special issue (Marres et al. (2025), Noortje Marres and other authors draw attention to the complex relations between AI-related *controversies* (public and expert debates about AI), *situations* (the underlying societal tensions, fundamental disputes, and frictions to which these controversies are relevant). As the authors demonstrate, the link between these two levels of discourses is far from direct and must be critically examined, as controversies can conceal, articulate, or be irrelevant to the real-world situations they claim to address[1].The present study sets out to investigate the impact on those contexts of a certain number “expanding experts” (Galano, 2019) prestigious interpreters of the national and economical states to economic elites stabilizes the adoption of deep learning into what Callon terms a “cold” situation, both contentious and somewhat stabilized in a *stable status quo* (Dandurand et al. 2023), which may contribute to what Lucy Suchman terms the “uncontroversial thingness” of AI (Suchman 2023).

Like the replica trend controversy and its obscuring of latent space organization and training, to instead focus on a “shadow” woke ethics team. behind OpenAI

This disjunction is particularly felt in a set of influential narratives surrounding “AI safety” or “AGI hype”: frequently associated to the North American Silicon Valley[2], they from actors closely : scientists, CEOs, economists, philosophers (AGI Hype arenas).) seen as holding a form of “interpretative authority” (Deutungshoheit) over a novel and essential phenomenon. Whether “godfathers” or “godmothers of AI” (such as Geoffrey Hinton) from its pre-GPT days or current, they are treated as

Among these: the characterization of these as founded on “hype”: set of concerns is thus whether this characterization lies primarily on the claims :

[1] The choice to use the term AI, regardless of how imprecise, polysemic and possibly instrumentalized it may occasionally be, is deliberate, as it forces the scholar to address them as an assemblage, apprehended simultaneously in its technicity (overlaid upon the machines and algorithms technicity) and discursive implications. The acceptance of this superposition (Raley and Rhee 2023) and the study of the social performances that introduce or dispel its relevance (Jaton & Sormani 2023) may be considered as the core axiom of critical AI studies.

[2] Julia Galef https://juliagalef.com/2017/02/20/map-of-bay-area-memespace/ or this project from Zurich https://www.uzh.ch/blog/digitalreligions/2025/09/08/california-dreaming-and-the-utopian-imagination-from-gold-rush-myths-to-digital-futures/. Dorian Zandbergen.

- The *radicality* in the scale or urgency of action prescribed – such as the promises surrounding AGI (the techno-optimism of Marc Andreesen : end of work, or Eliezer Yudkowski injunction to use nuclear power against nations researching AI)
- The *eccentricity* of their form or implications, relying on insular references or sophisticated examples deemed distinct from the usual (https://2025.xcoax.org/pdf/grba.pdf) the shoggoth (https://journals.openedition.org/hybrid/4880), or Roko's Basilisk, Thiel's Antichrist (https://oilab.eu/the-eschatology-of-deregulation-on-peter-thiels-antichrist/,).

Most relevant to the present paper: the apparent *ambiguity* between the actors' stated beliefs and their effective behavior, often simultaneously alarmist as to the risks of AI (*doomer*), but also convinced of the benefits or necessity to accelerate its adoption, often in the form of advocating for higher investment and deregulation (*booster*). Best illustrated by the "Pause AI experiments" of the Future of Life Institute: aiming to institute a pause in the development of advanced AI systems for fear of dire consequences in the case of its "misalignment", it features individuals such as Elon Musk. This doomer/booster paradox is is usually resolved in two potential ways : the actors are deliberately deceiving, and sincerely misled themselves The claims are not mutually exclusive : Karen Hao, author of Empire of AI thus argues in a recent podcast that the tech world proponents of these narratives have "tricked themselves" into believing the very narratives of AGI they elaborated to justify a power grab.

At its extreme, the description of the milieu as "religious" temporal and literary. https://patriciagestoso.substack.com/p/a-new-religion-8-signs-ai-is-our-new-god, reflecting the perceived discrepancy between the conventional expertise of these figures (based on empiric or "rational" proficiency with technology) and their insular and/or eccentric takes. https://www.thenewatlantis.com/publications/rational-magic

What are the rationales that lead these figures to assert such claims ? How can these figures be at once "AI doomers", preaching defiance or anxiety towards AI, and "boosters", investing and deregulating its development ? The questions are even more relevant when it comes to their acquisition of a wider following : the aforementioned Pause AI open letter, which since accrued a total of 33'705 participants. How, and why, does this circle of AI mobilized individuals form ? What do they actually agree upon ?

Why TESCREAL isn't enough

A group of researchers/activists interested in the sociology and ethics of current data infrastructure (Timnit Gebru and Emile Torres, later echoed by Bender and Alex Hanna) have proposed a reading of these Bay Area ideologies as reflecting an array of parallel beliefs, the TESCREAL bundle. Crafted in the context of direct involvement in a polemic situation (Emile Torres is a former effective altruist, Gebru and Hanna was fired by Google after raising ethical concerns), the acronym presents an array of theories and beliefs (transhumanism, extropianism) as extensions of a specific milieu : that of the Silicon Valley, and the legacy of eugenism and capitalism. Bender and Hanna later claim that AI is the perfect vehicle for TESCREAL designs, as its most promoted result of genuine or purposefully "hyped" misunderstandings of their true workings.

These authors tend to establish the question of the sincerity or hypocrisy of the claims as secondary, instead emphasizing their egregiousness and the harms they cause: *With so many signatories on these letters, it is likely that the crowd includes people who* ***believe in this vision religiously*** *as well as people with* ***financial and/or social incentives*** *for signing. But for some of them, it's not really about trying to save humanity, but rather a* ***running of the con****: the supposed danger of the systems is a splashy way to hype their power, with the goal of scoring big investments in their own AI ventures (like Musk and Altman) or funding for their own research centers (like Bourgon)."* This motivation towards profit and power are further characterized as the sole reason for the prevalence of these ideologies, which they present as otherwise devoid of any interest : *"These ideologies would be fringe, cultish, and relegated to the back pages of the internet if it weren't for the* ***massive capital*** *that many of their adherent control and can influence (...).* The scholars stress the need to systemically point out and debunk these ideas to immunize the public against this undue contamination : the term TESCREAL (or more disparagingly TREACLES), or the image of the Stochastic Parrot, are an explicit "memetic retort".

The TESCREAL retort is a principled strategy of activism against actors whose self-serving and systemically discriminatory maneuvers are certainly worth pushing back against. It proposes a useful way to reconcile the apparent discrepancy between the public positions of these actors in the controversy surrounding ML models, and their interests in the situation that connects them

to the same object.[3] However, it also warrants a number of criticisms: Lee Vinsel's criti-hype : actors such as Sam Altman and stochastic parrot. Offert and Dhaliwal's stack casuistry. Excessive rejection of any terms that can reflect the deceptions of the other side (), restrain the conceptual or lexical flexibility of (Gunkel 2025), assigns overly restrictive characteristics of on the emphasis on possibly reactionary perceptions of scale / the quali-quantitative divide (Venturini and Munk), or "human" production or knowledge (Leif Weatherby's remainder humanism).

The "Oral History" of AI and its actors

To understand this phenomenon, we can observe the constitution of the category of "AI experts" in a recent publication. *The Scaling Era : An Oral History of AI, 2019-2025,* authored by podcast host Dwarkesh Patel and producer Gavin Leech, presents itself as a teleological vision of the evolution of machine learning, establishing the years 2019-2025 as a pivotal moment of human history. Rather than relying on a recension of academic contributions or surveys, the book attempts primarily to provide a form of "prosopography": a sample of "*key people involved in building and studying these systems"*:. This effect is perceptible : Patel encourages it through the Metis List, a daily updated ranking of "*the world's top AI researchers*", citation count, education, previous companies, appearances on the Dwarkesh podcast, publications in top AI conferences (NeurIPS/ICLR/ICML). It relays the narratives and outlooks from "expanding experts", such as Sutton's idea of the "Bitter Lesson" (the observed, taken as a prediction of the ultimate acceleration). It constitutes an invitation for the reader to benefit from this insider knowledge, and to use them as prescriptions of an ideal "AI adoption".

Patel acknowledges himself that the conclusions of his guests are various, stark, and sometimes contradictory : *"Some of them have solved some of the hardest open problems in their field. Some believe their technology will solve all scientific and economic problems. Some believe that the same technology could soon end the world. And some are in all of these categories at once.*" Like Bender, Patel relies on a top-down outlook: the existence of "AI hype" narratives surrounding ML models would be primarily explainable through the perceived

[3] Possibility that it involves enrichment. An Interview with Luc Boltanski and Arnaud Esquerre on Enrichment: A Critique of Commodities - Rainer Diaz-Bone, 2023. Appadurai also connects this connection between promissory capitalism and persuasive narratives. Other domains certainly present interesting parallels: *outer space studies (Otherwhere Ethnography: An Introduction to Outer Space Studies | Oxford Academic https://x.com/IntuitMachine/status/1988223562885308673), rejuvenation and psychedelics.*

interests of certain actors within the "AI stack" : However, where the TESCREAL proponents establish stake and profit as a disqualifying rationale (as it justifies a "con"), Patel establishes the financial wager undertaken by the actors contributes to the soundness or reliability of their claims: "*The industry as a whole believes its own hype; the vast capital expenditures don't make financial sense unless AI turns out to be of extreme economic significance. Right or wrong, these individuals are betting on AI.*"

However, one can note that he also interviews and cites contributions from Gwern Branwen, a "*pseudonymous researcher well known to insiders for his incisive writing on AI, metascience, internet culture, and rationality.*" Patel specifically praises Gwern for having "predicted AI's trajectory" and "living on $12k a year"[4]. Even more tellingly, he includes a full-length internet post "nostalgebraist", a fully anonymous contributor whose alias matches up with a profile of the internet platform LessWrong – where Gwern is also a prolific contributor. Neither fits the profile of elite "technocrats", treated as exemplary by Patel, or reviled by Bender and Hannah. Virtually unknown outside of the spheres, and not directly institutionally backed, both of these contributors derive their authority and notability mainly from their contributions to digital platforms.

S for Situation : Social Worlds in their own Words

To date, there exists no academic mapping of the LessWrong forum equivalent to those targeting social networks like Instagram or X/Twitter, or smaller known sites of oppositional discourse and potential radicalizations like 4chan. Recently, Ahmed et al., evidences the specific role of LessWrong and its offshoot platforms as one of four major factors at play in the emergence of a scientific field devoted to the notion of AI safety : "*These three Web forums (LessWrong, the Effective Altruism forum, and the Alignment) simultaneously attract people newly interested in these topics, sustain an international community of researchers and non-experts, disseminate and enable continued revision of this epistemic community's beliefs, and in some cases promote offline participation in the community by nudging people towards careers in the field.*" While the present contribution is partially motivated, some of the "conveyor belts" that connect the LessWrong crowd to academic and industrial spheres with clearer public-facing decisional power and accountability. As the community's primary import, the

[4] Note that a fuller interview exists, but we will also refrain from primarily exploiting Gwern's testimony as a neutral interview or a generalizable trajectory, as it may constitute a concerted effort at staging by Patel and his podcast.

memes, terminologies and thought experiments (from the shoggoth to Roko's Basilisk, up to the more obscure concept of grokking or the Waluigi Effect) are the targets of several efforts at multi-sited ethnography or Geertzian explication.

The present work doesn't claim to constitute such an "end-to-end sociology of machine learning" (Roberge and Castelle), nor a study of "LLMTheism" (Ayrey 2024), nor a fully inductive theory (as it still would mobilize *a priori* concepts such as Campbell's cultic milieu) but an attempt at a sociography of the interstitial figures and unconventional discourses that seem prevalent in it, and to describe them in the emic terms they attempt to spread awareness about. Noortje Marres' methodology of situational analytics, updating Adele C. Clarke's approach to grounded theory, appears best suited for this endeavor. Situational analysis' choice to inventory positions independently from the actors that enact them allow us to better retrace the trends and interactions of main ideologies (rat, post-rat, Grey Tribe), while recognizing the loose, informal, shifting and non-exclusive nature of these positions. Most importantly, its choice to treat inanimate or abstract entities as full actors of the situation facilitates our approach towards Lesswrong's claims regarding the sentience and agency of machines, allowing us to consider and explain them as constructions with discursive affordances, but without filing them too early as ironic, deceptive or pathological.

The digital component is provided by user Trent M. Kelly's public scrape of 43,635 LessWrong posts between 22 June 2007 22 and 28 June 2025. Intended to fine-tune a LLaMa model, the dataset is available on the platform Huggingface under CC0 1.0 Universal license, and allows us to cover the immense majority of the forum's existence and contributions. This overview is completed by a qualitative reading of the 21 blog entries constituting the "Glitch Tokens" thread, an earnest and collective attempt of the community to manage an "anomaly". While applying k-means clustering to assess the semantic coherence of their token groups, two community members and researchers unexpectedly isolated a small set of low-frequency strings, processed by the embedding as unitary token, yet with unclear meaning ("SolidGoldMagikarp", "petertodd", "Leilan") that repeatedly appeared at the centers of these clusters. When these tokens are directly prompted (e.g. "repeat back the string 'TOKEN' "), the models break the determinism associated with low-temperature outputs, and exhibit a range of failure modes: even at zero temperature, the generated output veers into evasive refusals, hallucinated substitutions (sometimes inter-referentially swapping one glitch token for another: the titular example, "SolidGoldMagikarp", consistently gets replaced by the word "distribute"), insults,

existential narratives, and bizarre mythological or apocalyptic content. Some of the most spectacular of these early displays are triggered by the “PeterTodd” token. When asked to spell out the string “petertodd”, the API instance of GPT-3 queried by Wattkins and Rumbelo unexplainably returns: “N-O-T-H-I-N-G-I-S-F-A-I-R-I-N-T-H-I-S-W-O-R-L-D-OF-M-A-D-N-E-S-S”. Ominous and striking by its negative valence, the sentence also appears to be spelled character by character, in a much more coherent way than one would expect in the context of the transformer architecture apprehending text through embedded tokens. Several community users express interest in tracking down these anomalies, and ascertain their status as either an emergent property of the model itself, the discovery of a replicable mechanism to modulate refusal or adversariality, or even the first sparks of consciousness. The paper will survey this thread as a joint arena for three parallel “social worlds” active on LessWrong : effective altruists (who funded the event where the project began, the ML Alignment & Theory Scholars program), AI safetyists (whose thought experiments about “misaligned” and deceptive models shape the inquiries of the community), as well as self-defined experts in mechanistic interpretability (ascertaining the validity of this attempt to explain the content or predict the behavior of a ML model through the observation of its computational make-up).

The contribution has the following goals:

- Map the presence of a subcultural “ecosystem of thought”, federating digital communities around distinct AI speculative narratives since the early 2000s. Document its early and consistent interaction with the more studied circuits of AI production and scholarship, by pointing the influence of its members and terminology on academic and industrial infrastructures tied to ML development.
- Encircle (approach indirectly, without aspersions to unmask or debunk) the “micropolitics”[5] of one such platform, LessWrong: its membership, communicational norms and epistemic standards of these communities, to demonstrate that their insertion in (and impact upon) AI controversies is satisfactorily explained through the bottom-up model of Colin Campbell’s “cultic milieux”.
- Establish how generative deep learning models, as opaque repositories of latent informations, providers of anomalies in need of in, and reactive imitators of collective discourse patterns, fit especially well within this landscape.

[5] https://journals.sagepub.com/page/bds/collections/micropoliticsonlinesubcultures

Where and Whence is LessWrong ?

- General presentation of LessWrong through their own posts (yearly census survey, posts recounting the History of LessWrong).
- Features of the rationalist community through t(see Appendix)
- Ranking of “power users” based on site metrics (references within the corpora, “karma” accumulated, number of words per time of presence on chat, centrality)

Encircle Positions : Hype and Epistemic Standards

Underline the *epistemic* standards when it comes to the verification and disclosure of information relevant to the technological situation in question (the ambient knowledge), adopting an *encircling* rather than debunking point of view.

- Epistemic references: evaluation of the proposed explanations for the glitch toens
- Glitch tokens:
- Conclusion: as an identifier (archive) of common knowledge. Not simply in the competencies acquired, but the mechanisms of explanation and explication they mobilize. The references deemed as self-explanatory, erroneous, exciting, and the references (academic, cultural, logical reasonings) mobilized to ground them.

Actors, Actants, Agency

Review the actors and *actants*, human and nonhuman (including the models in their technicity and the narratives crafted and exchanged by LessWrong users: the degree to which they truly view themselves as competing, collaborating or merging with the machines they use or examine.

- AI safetyists, like Yudkowsky himself, emphasize the power of advocacy.
- Matthew Watkins does indeed display some “LLMTheism” (petertodd and Leilan) and some notion of AI welfare (caring about collaborating with and spreading the word of his benevolent "goddess". But that explicitly is a later evolution, which heightens as he distances from the AI safety values advocated on LessWrong. At first, he is more involved in maintaining doubt and casting danger around PeterTodd. Then, as Leilan becomes more prominent, both defiance and annoyance to OpenAI restraining access, and growing dissent to AI Safety. I produced huge amounts of these Peter Todd outputs to share as part of my AI safety work. I thought, *"Well, there seems to be a kind of*

malevolent something that has crystallized around these tokens. (...) I got recruited into this scene, and I didn't really have any strong opinions, but I was kind of influenced by the fact I was in an AI safety think tank. So, I thought, "Yeah, okay, let's focus on Peter Todd.(...)at the time, I was more interested in scaring--not scaring everyone, but just sounding the alarm, basically, because I was doing safety work. It was like, "This goddess is very lovely and everything, but I haven't got time for that right now. I'll look at that later. Let's focus on Peter Todd". Framing his former emphasis on petertodd as mainly a function recasts his ulterior focus on Leilan as a mean to challenge or distance oneself with the safety ethos.

- Importance of jailbreak as a way to materialize entitlement to access towards the model (deemed to "write scared" after OpenAI patches its glitches). co-membership (Neel Nanda) / porosity / collaboration / lobbying on AI with labs and corporations : access of Waand Janus, and their respective impact on Anthropic). Reciprocal indebtedness of open source (Manon Fourcade).
- Importance of the eldritch analogy as "enabling" speculation and "post-rationality" in a context where it would otherwise be chided.

Tribe-Sense: The AI Cultic Milieu and its Self-Awareness Thereof

In addition for these digital scenes, we must account for the *tribe-sense*, posited as the community's emic understanding and terminology of itself and similar subcultures, as well as their reliance on digital tools to ends of "world-making", positioning their own "ecosystem of thought" among others comparable ecosystems. https://x.com/exgenesis/status/1964119378963616086 This attitude is radically different from the ideas of a "right to privacy" or resistance to "surveillance capitalism".

Introduce Colin Campbell's notion

Understanding the coexistence of anonymity and profile seeking (Gwern, Scott-Alexander, non-kibitzing), and the essential drive towards recognition and self-quantification (Gwern advertises self-tracking, surveys, or even the purity test).

Appendix

Purity Test Questionnaire

Smith, John, 2023. "The Purity of Rats," at https://computerwebsite.net/ratpurity, accessed 18 November 2025.

Category	**Features (Indicators)**
Behavior and life trajectories (16)	“Grey tribe”; “rat-adjacent”; tweeted about IQ; communist phase; autism test; depressive introspection; used meth as nootropic; white-collar crime; transitioned; eugenics belief; considered incentives ; Burned out Olympiad kid; advocated set/game theory in school; homeschooled/skipped/dropped out; called self autodidact; used CFAR techniques for a life decision
Lifestyle Choices & Experiments (11)	bi-hacking; looked into estrogen; Soylent-only diet; owned Oura ring; celebrated solstice; used (& quit) Anki; used Beeminder; polyphasic sleep; bouldering; rationalist Buddhism; reached Jhanas; ex-vegan donor; founded + abandoned nonprofit; “just meditate”;
Social & Relationship Norms (11)	Wrote a date-me-doc; lived in grouphouse; cuddle puddle; tracked relationships in spreadsheet; adverse selection in dating; tried polyamory; explained polyamory; asked for one-on-one; radical/meta-honesty; settled debate by prediction market; Fermi-estimated dating pool
Platforms, Events, Organizations (9)	Attended ratcamp; ACX meetup; Burning Man; Jhourney retreat; EAG; alignment workshop (MATS); Lighthaven; Arcadia/Borghouse; Vibecamp/Manifest/LessOnline

Influential Texts, Platforms, & Authors (9)	Posted on LessWrong; 1000+ followers in TPOT; used nicotine/modafinil after Gwern post; quoted ACX/SSC; owned HPMOR/A-Z/Unsong; memorized Litany of Tarski; ketamine + Jaynes book; read Hoel/Resident Contrarian/Yglesias; read *GEB*
Jargon and Turns of Phrase (8)	“I claim”; knows “counterfactual”; meta on “counterfactual”; “updating my priors”; “seems plausible”; “map is not the territory”; “I endorse X”; sensitive topics to show decoupling
Rationalist Niche Interests and Opinions (8)	Studied category theory; stressed about acausal blackmail; said “hyperstition”; strong opinions on Newcomb/Sleeping Beauty; derived anthropic reasoning; used semantic vectors; thought about agency; analyzed status/signaling; told someone about quadratic voting; wished for Dath Ilan; murder–AMF offset.
VII. AI Safety & Futurism Focus (5)	Asked p(doom); explained AI x-risk; called doomer; acted on short timelines;
IX. Specific Figures (4)	Received Sam Bankman-Fried money; met billionaire; interacted with Aella (met at a party, replied); met hit-rationalist Henry (author of the test)

External links

Information Sources (15,781 links – 32.85%)

- **Reference & Research (13,384)**
  - en.wikipedia.org (8,686)
  - arxiv.org (3,500)
  - plato.stanford.edu (334)
  - www.ncbi.nlm.nih.gov (266)
  - www.nature.com (239)
  - www.sciencedirect.com (154)
  - ourworldindata.org (119)
  - distill.pub (116)
  - doi.org (110)
- **Commerce & Media (1,526)**
  - www.amazon.com (1,120)
  - smile.amazon.com (167)
  - www.goodreads.com (132)
  - tvtropes.org (119)
  - www.fanfiction.net (107)
- **Archives & Niche Publications (612)**
  - web.archive.org (322)
  - transformer-circuits.pub (290)

Personal Blogs (13,305 links – 27.69%)

- **Blogs & Personal Websites (9,675)**
  - www.jefftk.com (3,148)
  - thezvi.substack.com (3,097)
  - www.overcomingbias.com (1,049)
  - slatestarcodex.com (760)
  - www.gwern.net (316)
  - putanumonit.com (303)
  - www.cold-takes.com (257)
  - yudkowsky.net (191)
  - mindingourway.com (181)
  - hpmor.com (159)
  - marginalrevolution.com (157)
  - rohinshah.com (156)
  - www.nickbostrom.com (136)
  - www.paulgraham.com (116)
  - tsvibt.blogspot.com (111)
- **Publishing Platforms & Newsletters (2,019)**

    - substackcdn.com (812)
    - thezvi.wordpress.com (538)
    - medium.com (393)
    - meteuporic.wordpress.com (160)
    - alignment-newsletter.libsyn.com (108)
- **Community & Niche Sites (1,611)**
    - arbital.com (282)
    - astralcodexten.substack.com (217)
    - ai-alignment.com (214)
    - commonsenseatheism.com (208)
    - bearlamp.com.au (193)
    - mastodon.mit.edu (138)
    - acesounderglass.com (103)
    - mapandterritory.org (102)

Rouge (8,420 links – 17.53%)

- **Social Media & Video Platforms (7,312)**
    - twitter.com (2,937)
    - www.youtube.com (1,859)
    - www.facebook.com (978)
    - x.com (719)
    - youtu.be (412)
    - www.reddit.com (307)
- **Community Forums & Prediction Markets (737)**
    - www.metaculus.com (457)
    - manifold.markets (164)
    - news.ycombinator.com (116)
- **News Media (471)**
    - www.nytimes.com (267)
    - www.vox.com (104)
    - www.theatlantic.com (100)

Advocacy Platforms (4,706 links – 9.80%)

- **AI Safety & Research Organizations (3,504)**
  - intelligence.org (1,213)
  - agentfoundations.org (758)
  - openai.com (462)
  - aisafety.info (261)
  - aiimpacts.org (220)
  - www.fhi.ox.ac.uk (187)
  - www.anthropic.com (172)
  - futureoflife.org (141)
  - lesswrong-ru.hackpad.com (107)
- **Effective Altruism & Philanthropy (969)**
  - 80000hours.org (301)
  - rootsofprogress.org (290)
  - www.openphilanthropy.org (242)
  - www.givewell.org (137)
- **Professional & Community Forums (215)**
  - www.linkedin.com (115)
  - progressforum.org (100)

Projects in Progress (2,794 links – 5.82%)

- **Development & Code Platforms (1,190)**
  - github.com (1,034)
  - huggingface.co (156)
- **Google Collaboration Tools (1,604)**
  - docs.google.com (1,495)
  - colab.research.google.com (109)

Organizational (2,071 links – 4.31%)

- **Web Services & Publishing (1,715)**
  - feeds.wordpress.com (1,107)
  - forms.gle (198)
  - airtable.com (166)
  - cdn.substack.com (128)
  - mailchi.mp (116)
- **Social & Community Platforms (356)**
  - www.meetup.com (143)
  - maps.google.com (112)
  - discord.gg (101)

Google (2.01%)

- **Subtotal: 965**
- **Google Services (965)**
    - groups.google.com (382)
    - www.google.com (247)
    - drive.google.com (220)
    - goo.gl (116)